Yulin Wang

Ph.D. Student, Tsinghua University

☎ +86 131-2121-5969   ●   ✉ wang-yl19@mails.tsinghua.edu.cn

🌐 www.wyl.cool

Education

Department of Automation, Tsinghua University

Ph.D. Student in Machine Learning and Computer Vision   *2019.09 – 2025.04 (expected)*

- Advisors: Prof. Cheng Wu and Prof. Gao Huang.

School of Automation Science and Electrical Engineering, Beihang University

B.Eng. in Automation   *2015.09 – 2019.06*

- GPA **Top 1/231**.
- Recipient of the **"Shen Yuan" Medal** (the highest honor for undergraduate students, **10/18,000+**).

Research Experience

Berkeley Deep Drive, University of California, Berkeley

Research Intern (advisor: Dr. Ching-Yao Chan)   *2018.07 – 2018.09*

Lab of Intelligent Manufacturing, Beihang University

Research Intern (advisor: Prof. Fei Tao)   *2017.06 – 2018.06*

Research Interests

Yulin Wang's research has centered on addressing challenges in *computational efficiency* and *data efficiency* when building and deploying large-scale foundation AI models (*e.g.*, large vision/multi-modal models, generative models, and embodied foundation models for robotics). His goal is to enable high-level intelligence at an affordable training and inference cost.

During his Ph.D., he discovered that introducing human-like *adaptiveness* into deep learning could be an effective approach to achieving this goal. Specifically, he took inspiration from human cognition processes and neuroscience, and developed algorithms capable of *adaptively adjusting inference and learning strategies conditioned on the diversified information and patterns within different data*.

Publications

Yulin Wang has published a number of works in top-tier conferences & journals in the fields of machine learning and computer vision, including *TPAMI* (**6**), *IJCV* (**1**), **NeurIPS** (**5**), **ICLR** (**1**), **ICCV** (**6**), **CVPR** (**5**), and **ECCV** (**2**). He has collected **more than *2,400 citations*** according to **Google Scholar**.

I. Selected Publications

I.1 Adaptive inference for large-scale perception/generative models

> **TL;DR:** *Dynamically adjusting inference strategies based on each of the diversified input/generated data samples (e.g., mimicking human vision, modeling visual perception as a coarse-to-fine sequential decision process with spatial-temporal dynamic computation), thus improving computational inference efficiency.*

[1] - **Yulin Wang**, Haoji Zhang, Yang Yue, Shiji Song, Chao Deng, Junlan Feng, Gao Huang
   Uni-AdaFocus: Spatial-temporal Dynamic Computation for Video Recognition
   IEEE Transactions on Pattern Analysis and Machine Intelligence (TPAMI, IF$_{5\text{-year}}$=22.2), 2024

- **Yulin Wang**, Zhaoxi Chen, Haojun Jiang, Shiji Song, Yizeng Han, Gao Huang
   Adaptive Focus for Efficient Video Recognition
   IEEE/CVF International Conference on Computer Vision (ICCV Oral, acceptance rate: 3%), 2021

[2] - Gao Huang*, **Yulin Wang***, Kangchen Lv, Haojun Jiang, Wenhui Huang, Pengfei Qi, Shiji Song
      [*co-first author with my advisor*]
      Glance and Focus Networks for Dynamic Visual Recognition
      IEEE Transactions on Pattern Analysis and Machine Intelligence (*TPAMI, IF*$_{5\text{-year}}$**=22.2**), 2023
    - **Yulin Wang**, Kangchen Lv, Rui Huang, Shiji Song, Le Yang, Gao Huang
      Glance and Focus: A Dynamic Approach to Reducing Spatial Redundancy in Image Classification
      Advances in Neural Information Processing Systems (**NeurIPS**), 2020

[3] **Yulin Wang**, Rui Huang, Shiji Song, Zeyi Huang, Gao Huang
      Not All Images are Worth 16x16 Words: Dynamic Transformers for Efficient Image Recognition
      Advances in Neural Information Processing Systems (**NeurIPS**), 2021

I.2 Efficient training of foundation models

TL;DR: Designing tailored learning strategies for each type of discriminative pattern within data (e.g., low/high-frequency patterns and local/global patterns in vision), and organizing them into a unified training procedure mimicking human learning (e.g., through curriculum learning or neuroscience-inspired mechanisms), thus training large foundation models more efficiently, stably, and effectively.

[4] - **Yulin Wang**, Yang Yue, Rui Lu, Yizeng Han, Shiji Song, Gao Huang
      EfficientTrain++: Generalized Curriculum Learning for Efficient Visual Backbone Training
      IEEE Transactions on Pattern Analysis and Machine Intelligence (*TPAMI, IF*$_{5\text{-year}}$**=22.2**), 2024
    - **Yulin Wang**, Yang Yue, Rui Lu, Tianjiao Liu, Zhao Zhong, Shiji Song, Gao Huang
      EfficientTrain: Exploring Generalized Curriculum Learning for Training Visual Backbones
      IEEE/CVF International Conference on Computer Vision (**ICCV**), 2023

[5] **Yulin Wang**, Zanlin Ni, Shiji Song, Le Yang, Gao Huang
      Revisiting Locally Supervised Learning: An Alternative to End-to-end Training
      International Conference on Learning Representations (**ICLR**), 2021

I.3 Data-efficient learning/fine-tuning

TL;DR: Mimicking "learning by analogy", augmenting the semantic diversity of training samples based on their individual characteristics, thus reducing the expensive cost of collecting high-quality training data or fine-tuning data (e.g., for supervised fine-tuning and transfer learning).

[6] - **Yulin Wang**, Gao Huang, Shiji Song, Xuran Pan, Yitong Xia, Cheng Wu
      Regularizing Deep Networks with Semantic Data Augmentation
      IEEE Transactions on Pattern Analysis and Machine Intelligence (*TPAMI, IF*$_{5\text{-year}}$**=22.2**), 2022
    - **Yulin Wang**, Xuran Pan, Shiji Song, Hong Zhang, Cheng Wu, Gao Huang
      Implicit Semantic Data Augmentation for Deep Networks
      Advances in Neural Information Processing Systems (**NeurIPS**), 2019

II. All Publications

II.1 All Publications 1/3 – Adaptive inference for large-scale perception/generative models

[7] **Yulin Wang**, Yang Yue, Xinhong Xu, Ali Hassani, Victor Kulikov, Nikita Orlov, Shiji Song, Humphrey Shi, Gao Huang
      AdaFocus V3: On Unified Spatial-temporal Dynamic Video Recognition
      European Conference on Computer Vision (**ECCV**), 2022

[8] **Yulin Wang**, Yang Yue, Yuanze Lin, Haojun Jiang, Zihang Lai, Victor Kulikov, Nikita Orlov, Humphrey Shi, Gao Huang
AdaFocus V2: End-to-End Training of Spatial Dynamic Networks for Video Recognition
*IEEE/CVF Conference on Computer Vision and Pattern Recognition (**CVPR**), 2022*

[9] Yang Yue*, **Yulin Wang***, Bingyi Kang, Yizeng Han, Shenzhi Wang, Shiji Song, Jiashi Feng, Gao Huang   [*co-first author, supervising Yang Yue (junior Ph.D. student)*]
Dynamic Inference of Multimodal Large Language Models for Efficient Robot Execution
*Advances in Neural Information Processing Systems (**NeurIPS**), 2024*

[10] Zanlin Ni*, **Yulin Wang***, Renping Zhou, Yizeng Han, Jiayi Guo, Zhiyuan Liu, Yuan Yao, Gao Huang   [*co-first author, supervising Zanlin Ni (junior Ph.D. student)*]
ENAT: Rethinking Spatial-temporal Interactions in Token-based Image Synthesis
*Advances in Neural Information Processing Systems (**NeurIPS**), 2024*

[11] Zanlin Ni*, **Yulin Wang***, Renping Zhou, Rui Lu, Jiayi Guo, Jinyi Hu, Zhiyuan Liu, Yuan Yao, Gao Huang   [*co-first author, supervising Zanlin Ni (junior Ph.D. student)*]
AdaNAT: Exploring Adaptive Policy for Token-Based Image Generation
*European Conference on Computer Vision (**ECCV**), 2024*

[12] Zanlin Ni*, **Yulin Wang***, Renping Zhou, Jiayi Guo, Jinyi Hu, Zhiyuan Liu, Shiji Song, Yuan Yao, Gao Huang   [*co-first author, supervising Zanlin Ni (junior Ph.D. student)*]
Revisiting Non-Autoregressive Transformers for Efficient Image Synthesis
*IEEE/CVF Conference on Computer Vision and Pattern Recognition (**CVPR**), 2024*

[13] Ziwei Zheng, Le Yang, **Yulin Wang**, Miao Zhang, Lijun He, Gao Huang, Fan Li
Dynamic Spatial Focus for Efficient Compressed Video Action Recognition
*IEEE Transactions on Circuits and Systems for Video Technology (**TCSVT, IF**$_{5\text{-year}}$**=7.1**), 2023*

[14] Yizeng Han, Gao Huang, Shiji Song, Le Yang, Honghui Wang, **Yulin Wang**
Dynamic Neural Networks: A Survey
*IEEE Transactions on Pattern Analysis and Machine Intelligence (**TPAMI, IF**$_{5\text{-year}}$**=22.2**), 2022*

[15] Ziwei Zheng, Zechuan Zhang, **Yulin Wang**, Shiji Song, Gao Huang, Le Yang
Rethinking the Architecture Design for Efficient Generic Event Boundary Detection
*ACM International Conference on Multimedia (**ACM MM**), 2024*

[16] Yizeng Han, Dongchen Han, Zeyu Liu, **Yulin Wang**, Xuran Pan, Yifan Pu, Chao Deng, Junlan Feng, Shiji Song, Gao Huang
Dynamic Perceiver for Efficient Visual Recognition
*IEEE/CVF International Conference on Computer Vision (**ICCV**), 2023*

[17] Yifan Pu, Yiru Wang, Zhuofan Xia, Yizeng Han, **Yulin Wang**, Weihao Gan, Zidong Wang, Shiji Song, Gao Huang
Adaptive Rotated Convolution for Rotated Object Detection
*IEEE/CVF International Conference on Computer Vision (**ICCV**), 2023*

[18] Le Yang, Haojun Jiang, Ruojin Cai, **Yulin Wang**, Shiji Song, Gao Huang, Qi Tian
CondenseNet V2: Sparse Feature Reactivation for Deep Networks
*IEEE/CVF Conference on Computer Vision and Pattern Recognition (**CVPR**), 2021*

II.2 All Publications 2/3 – Efficient training of foundation models

[19] Zanlin Ni*, **Yulin Wang***, Jiangwei Yu, Haojun Jiang, Yue Cao, Gao Huang
[*co-first author, supervising Zanlin Ni (junior Ph.D. student)]
Deep Incubation: Training Large Models by Divide-and-Conquering
*IEEE/CVF International Conference on Computer Vision (**ICCV**), 2023*

II.3 All Publications 3/3 – Data-efficient learning/fine-tuning

[20] Chaoqun Du, **Yulin Wang**, Shiji Song, Gao Huang
Probabilistic Contrastive Learning for Long-Tailed Visual Recognition
*IEEE Transactions on Pattern Analysis and Machine Intelligence (**TPAMI**, **IF**$_{\text{5-year}}$**=22.2**), 2024*

[21] - Mixue Xie, Shuang Li, Kaixiong Gong, **Yulin Wang**, Gao Huang
Adapting Across Domains via Target-Oriented Transferable Semantic Augmentation Under Prototype Constraint
*International Journal of Computer Vision (**IJCV**, **IF**$_{\text{5-year}}$**=14.5**), 2023*
- Shuang Li, Mixue Xie, Kaixiong Gong, Chi Harold Liu, **Yulin Wang**, Wei Li
Transferable Semantic Augmentation for Domain Adaptation
*IEEE/CVF Conference on Computer Vision and Pattern Recognition (**CVPR Oral**), 2021*

[22] Yifan Pu, Yizeng Han, **Yulin Wang**, Junlan Feng, Chao Deng, Gao Huang
Fine-grained Recognition with Learnable Semantic Data Augmentation
*IEEE Transactions on Image Processing (**TIP**, **IF**$_{\text{5-year}}$**=12.1**), 2023*

[23] Wenxuan Ma, Shuang Li, Jinming Zhang, Chi Harold Liu, Jingxuan Kang, **Yulin Wang**, Gao Huang
Borrowing Knowledge From Pre-trained Language Model: A New Data-efficient Visual Learning Paradigm
*IEEE/CVF International Conference on Computer Vision (**ICCV**), 2023*

[24] Wenxuan Ma, Jinming Zhang, Shuang Li, Chi Harold Liu, **Yulin Wang**, Wei Li
Making the Best of Both Worlds: A Domain-Oriented Transformer for Unsupervised Domain Adaptation
*ACM International Conference on Multimedia (**ACM MM**), 2022*

[25] Shuang Li, Kaixiong Gong, Chi Harold Liu, **Yulin Wang**, Feng Qiao, Xinjing Cheng
MetaSAug: Meta Semantic Augmentation for Long-Tailed Visual Recognition
*IEEE/CVF Conference on Computer Vision and Pattern Recognition (**CVPR**), 2021*

Selected Awards and Honors

○ National Scholarship for Ph.D. Students, Ministry of Education of China · · · · · · · · · · · *2021, 2023*
  (2 times, Top 2% Ph.D. students in Tsinghua University)

○ **Microsoft Research Ph.D. Fellowship**, Microsoft Research Asia · *2022*
  (12 Ph.D. students in the Asia-Pacific region)

○ **ByteDance Scholarship**, ByteDance Ltd. · *2022*
  (10 Ph.D. students in China)

○ **Baidu Scholarship**, Baidu Inc. · *2021*
  (10 Ph.D. students worldwide)

○ **CCF-CV Outstanding Young Researcher Award**, China Computer Federation (CCF) · · · *2021*
  (3 Ph.D./MS students in China)

- Travel Award, NeurIPS *2019*
- **"Shen Yuan" Medal**, Beihang University *2018*
  (the highest honor for undergraduate students, 10/18,000+)
- National Scholarship for Undergraduates, Ministry of Education of China *2017, 2018*
  (**2 times, Top 2% undergraduates in Beihang University**)
- "Gong Xin" Innovation Scholarship, Ministry of Industry and Information Technology of China *2017*
  (**Top 1/231 in Beihang University**)
- First Prize, National Undergraduate Mathematical Contest in Modeling *2017*
  (**Top 0.2%**)

Teaching Experience

- Guest Lecturer for *Pattern Recognition and Machine Learning* *Spring 2023, 2024*
- Teaching Assistant for *Nonlinear Programming* *Fall 2021, 2022, 2023*

Academic Service

- Reviewer for TPAMI, IJCV, TCYB, TNNLS, TCSVT, Pattern Recognition, TMLR, ...
- Reviewer for ICML, NeurIPS, ICLR, CVPR, ICCV, ECCV, AAAI, ...
  - *Outstanding Reviewer*, CVPR, 2021
- Co-sponsor of the Special Interest Group on Dynamic Neural Networks, Beijing Academy of Artificial Intelligence (BAAI).
  - https://littlepure2333.github.io/dynamic-neural-network
  - Core members include more than 20 researchers from 8 universities. We have organized more than 30 academic reports and tutorials. The cumulative audience has exceeded 1,000.

Invited Talks and Presentations

- 2023.09, Tsinghua-Berkeley Shenzhen Institute, Tsinghua University, Dynamic Inference of Neural Networks
- 2023.02, School of Automation, Beijing Institute of Technology, Vision Transformers Meet Dynamic Inference
- 2021.12, PRCV 2021, Dynamic Deep Networks for Reducing Spatial Redundancy
- 2021.10, School of Computer Science, Fudan University, Dynamic Deep Networks for Reducing Spatial Redundancy
- 2021.09, Aibee (invited by Yuanqing Lin), Semantic Data Augmentation
- 2021.06, AI Time, Locally Supervised Deep Learning
- 2021.04, Beijing Academy of Artificial Intelligence, Dynamic Image/Video Recognition
- 2021.03, ByteDance Ltd., Semantic Data Augmentation
- 2020.11, Qingyuan Seminar, Glance and Focus Networks
- 2020.06, Huawei Technologies Ltd., Glance and Focus Networks
- 2019.10, School of Computer Science and Engineering, Beihang University, Semantic Data Augmentation